

Synchronization in Java SE 6 (HotSpot)

Dave Dice

`dice@sun.com`

HotSpot JVM Core Engineering

Synchronization Performance

- Contended costs (scalability + latency)
 - > Context switching is extremely expensive
 - > Unbounded spinning is unacceptable
 - > Address via **adaptive spinning**
- Uncontended costs (latency)
 - > Atomic CAS has high **local** latency
 - > 100s-1000s of cycles
 - > Address via
 - > **Biased Locking**
 - > **Lock Coarsening**
 - > Lock Elision through **Escape Analysis**

HotSpot Locking Fundamentals

- Object header - metadata
 - > Mark word
 - > Class pointer
 - > ... followed by constituent fields
- Mark word multiplexed
 - > Identity hashCode
 - > GC Age bits
 - > Synchronization information
 - > **Displaced** mark word

Object States – Encoded in Mark Word

- Neutral: Unlocked
- Biased: Locked|Unlocked + Unshared
 - > Tantamount to deferring unlock until contention
 - > Avoids CAS atomic latency in common case
 - > 2nd thread must **revoke** bias from 1st
- Stack-Locked: Locked + Shared but uncontended
 - > Mark points to displaced header on owner's stack
- Inflated: Locked|Unlocked + Shared and contended
 - > threads are blocked: enter or wait
 - > Mark points to heavy-weight objectmonitor structure

Key Observations

- Most objects are never locked
- If an object is locked it is usually locked by at most one thread during its lifetime
 - > Very few objects are locked by more than one thread
- Even fewer objects encounter contention
- Object type and allocation site correlate strongly with future synchronization behavior

Biased Locking

- Leverages the observation that most objects are locked by at most one thread in their lifetime
- *Bias* object O toward Thread T1
- T1 can then preferentially lock and unlock O without expensive atomic instructions (CAS)
- If T2 attempts to lock O we *revoke* bias from T1
 - > Either rebias to T2 or revert to normal locking and make O ineligible for further biased locking

Adaptive Spinning

- Spin-then-block strategy
 - > Try to avoid context switch by spinning on MP systems
- Spin duration
 - > Maintained per-monitor
 - > varies based on recent history of spin success/failure ratio
- Adapts to system load, parallelism, application modality
- MP-polite spinning
- Avoid spinning in futile conditions (owner is blocked)

HotSpot Locking Fundamentals (2)

- Fast-path cases inlined by JIT at synchronization site
- Revert to slow-path (native C code) when we need to park or unpark thread
- Platform-specific park-unpark to block and wake threads
- Slow-path monitor code is platform-independent
- Much faster than native mutex constructs for contended & uncontended cases (T2, windows)

Detecting Contention

- IDEs, Profilers or 3rd party tools
- Mpstat on Solaris – vctx rate
- If suspected, sample process with pstack
 - > Look near top of stack for threads blocked in **monitorenter** operations
- JVMStat (jstat) counters
 - > `jstat -J-Djstat.showUnsupported=true -snap <pid> | grep _sync_`

Detecting Contention (2)

- Dtrace:
 - > kernel “sched” provider
 - > hotspot-specific probes (*Recommended!*)
- Identify hot locks and break up into finer-grained locking
- Beware: adding more threads can sometimes reduce performance – application specific
 - > Particularly on Niagara
 - > Amdahl’s speedup law – parallel corollary
 - > Communication overhead can overwhelm parallelism benefit

New in 1.6

- No atomic/fence in common-case inline inflated exit path
- Code restructuring:
 - > Platform independent monitor code calls ...
 - > Platform-specific park-unpark
- Reduce futile wakeups
 - > Don't wake a thread in exit if thread woken in prior exit hasn't yet run
- Lock-free EntryList
- Adaptive spinning

New in 1.6 (2)

- Notify() moves thread from WaitSet to EntryList
 - > Previous versions actually woke notifyee
 - > Notifyee would simply jam on lock held by notifier
- Fairness vs throughput
 - > Optimized for system-wide throughput at the expense of short-term thread-specific fairness
 - > Succession policy: try to wake recently run threads
 - > Improved \$ and TLB utilization
- Better JSR166 (`java.util.concurrent`) support

New in 1.6 (3)

- Small changes to comply with JSR133
 - > Java Memory Model (JMM)
 - > JLS 3e, Chapter 17
 - > -XX:-UseBiasedLocking
- Biased Locking on by default
- Lock Coarsening on by default
 - > -XX:-EliminateLocks
- Lock Elision via Escape Analysis
 - > -XX:+DoEscapeAnalysis

1.6 Source Roadmap

- Slow-path native
 - > Platform-independent : Synchronizer.cpp
 - > Platform-specific park-unpark : os_<plaf>.cpp + .hpp
- Fast-path inlined
 - > Degenerate form of slow-path code
 - > C2 **FastLock** node
 - > assembler_sparc.cpp compiler_lock_object()
 - > <cpu>.ad for other architectures
 - > C1 c1_CodeStubs_<cpu>.cpp
 - > Template interpreter

Additional Information

- <http://blogs.sun.com/dave>



Synchronization in Java SE 6 (HotSpot)

Dave Dice

dice@sun.com

HotSpot JVM Core Engineering

Synchronization Performance

- Contended costs (scalability + latency)
 - > Context switching is extremely expensive
 - > Unbounded spinning is unacceptable
 - > Address via **adaptive spinning**
- Uncontended costs (latency)
 - > Atomic CAS has high **local** latency
 - > 100s-1000s of cycles
 - > Address via
 - > **Biased Locking**
 - > **Lock Coarsening**
 - > Lock Elision through **Escape Analysis**

Java SE 6
Copyright 2006 Sun Microsystems, Inc. All Rights Reserved.

HotSpot Locking Fundamentals

- Object header - metadata
 - > Mark word
 - > Class pointer
 - > ... followed by constituent fields
- Mark word multiplexed
 - > Identity hashCode
 - > GC Age bits
 - > Synchronization information
 - > **Displaced** mark word

Java SE 6
Copyright 2006 Sun Microsystems, Inc. All Rights Reserved.

Object States – Encoded in Mark Word

- Neutral: Unlocked
- Biased: Locked|Unlocked + Unshared
 - > Tantamount to deferring unlock until contention
 - > Avoids CAS atomic latency in common case
 - > 2nd thread must **revoke** bias from 1st
- Stack-Locked: Locked + Shared but uncontended
 - > Mark points to displaced header on owner's stack
- Inflated: Locked|Unlocked + Shared and contended
 - > threads are blocked: enter or wait
 - > Mark points to heavy-weight objectmonitor structure

Java SE 6
Copyright 2006 Sun Microsystems, Inc. All Rights Reserved.

Key Observations

- Most objects are never locked
- If an object is locked it is usually locked by at most one thread during its lifetime
 - > Very few objects are locked by more than one thread
- Even fewer objects encounter contention
- Object type and allocation site correlate strongly with future synchronization behavior

Java SE 6
Copyright 2006 Sun Microsystems, Inc. All Rights Reserved.

Biased Locking

- Leverages the observation that most objects are locked by at most one thread in their lifetime
- *Bias* object O toward Thread T1
- T1 can then preferentially lock and unlock O without expensive atomic instructions (CAS)
- If T2 attempts to lock O we *revoke* bias from T1
 - > Either rebias to T2 or revert to normal locking and make O ineligible for further biased locking

Java SE 6
Copyright 2006 Sun Microsystems, Inc. All Rights Reserved.

Adaptive Spinning

- Spin-then-block strategy
 - > Try to avoid context switch by spinning on MP systems
- Spin duration
 - > Maintained per-monitor
 - > varies based on recent history of spin success/failure ratio
- Adapts to system load, parallelism, application modality
- MP-polite spinning
- Avoid spinning in futile conditions (owner is blocked)

Java SE 6
Copyright 2006 Sun Microsystems, Inc. All Rights Reserved.

HotSpot Locking Fundamentals (2)

- Fast-path cases inlined by JIT at synchronization site
- Revert to slow-path (native C code) when we need to park or unpark thread
- Platform-specific park-unpark to block and wake threads
- Slow-path monitor code is platform-independent
- Much faster than native mutex constructs for contended & uncontended cases (T2, windows)

Java SE 6
Copyright 2006 Sun Microsystems, Inc. All Rights Reserved.

Detecting Contention

- IDEs, Profilers or 3rd party tools
- Mpstat on Solaris – vctx rate
- If suspected, sample process with pstack
 - > Look near top of stack for threads blocked in **monitorenter** operations
- JVMStat (jstat) counters
 - > `jstat -J-Djstat.showUnsupported=true -snap <pid> | grep _sync_`

Detecting Contention (2)

- Dtrace:
 - > kernel “sched” provider
 - > hotspot-specific probes (*Recommended!*)
- Identify hot locks and break up into finer-grained locking
- Beware: adding more threads can sometimes reduce performance – application specific
 - > Particularly on Niagara
 - > Amdahl’s speedup law – parallel corollary
 - > Communication overhead can overwhelm parallelism benefit

Java SE 6
Copyright 2006 Sun Microsystems, Inc. All Rights Reserved.

New in 1.6

- No atomic/fence in common-case inline inflated exit path
- Code restructuring:
 - > Platform independent monitor code calls ...
 - > Platform-specific park-unpark
- Reduce futile wakeups
 - > Don't wake a thread in exit if thread woken in prior exit hasn't yet run
- Lock-free EntryList
- Adaptive spinning

Java SE 6
Copyright 2006 Sun Microsystems, Inc. All Rights Reserved.

New in 1.6 (2)

- Notify() moves thread from WaitSet to EntryList
 - > Previous versions actually woke notifyee
 - > Notifyee would simply jam on lock held by notifier
- Fairness vs throughput
 - > Optimized for system-wide throughput at the expense of short-term thread-specific fairness
 - > Succession policy: try to wake recently run threads
 - > Improved \$ and TLB utilization
- Better JSR166 (java.util.concurrent) support

Java SE 6
Copyright 2006 Sun Microsystems, Inc. All Rights Reserved.

New in 1.6 (3)

- Small changes to comply with JSR133
 - > Java Memory Model (JMM)
 - > JLS 3e, Chapter 17
 - > -XX:-UseBiasedLocking
- Biased Locking on by default
- Lock Coarsening on by default
 - > -XX:-EliminateLocks
- Lock Elision via Escape Analysis
 - > -XX:+DoEscapeAnalysis

Java SE 6
Copyright 2006 Sun Microsystems, Inc. All Rights Reserved.

1.6 Source Roadmap

- Slow-path native
 - > Platform-independent : Synchronizer.cpp
 - > Platform-specific park-unpark : os_<plaf>.cpp + .hpp
- Fast-path inlined
 - > Degenerate form of slow-path code
 - > C2 **FastLock** node
 - > assembler_sparc.cpp compiler_lock_object()
 - > <cpu>.ad for other architectures
 - > C1 c1_CodeStubs_<cpu>.cpp
 - > Template interpreter

Java SE 6
Copyright 2006 Sun Microsystems, Inc. All Rights Reserved.

Additional Information

- <http://blogs.sun.com/dave>

Java SE 6

Copyright 2006 Sun Microsystems, Inc. All Rights Reserved.